

Started on	Thursday, 30 January 2025, 11:55 AM
State	Finished
Completed on	Thursday, 30 January 2025, 12:00 PM
Time taken	4 mins 59 secs
Marks	10.00/10.00
Grade	100.00 out of 100.00

Question 1

Complete

Mark 1.00 out of 1.00

What is the size or shape of Q, K, and V in the ScaledDotProduct Attention forward pass ? (Assuming batch size is B, the sequence length is L, and the hidden size is H)

- ☐ a. (H, B, L)
- ☒ b. (B, L, H)
- ☐ c. (B, H, L)
- ☐ d. (L, B, H)

Question 2

Complete

Mark 1.00 out of 1.00

What does the term "scaled dot-product attention" refer to?

- ☐ a. Scaling the dot product of query and key vectors by the softmax operation
- ☐ b. Applying batch normalization to the query and key vectors
- ☐ c. Computing attention weights using convolution operations
- ☒ d. Scaling the dot product of query and key vectors by the square root of the key vector dimension

Question 3

Complete

Mark 1.00 out of 1.00

What role does the "key" vector play in the attention mechanism?

- ☐ a. It is the final output of the attention module
- ☐ b. It is an initialization for positional embeddings
- ☒ c. It is used to compute the similarity with the query
- ☐ d. It determines the scaling factor for normalization

**Question 4**

Complete

Mark 1.00 out of 1.00

What are the three main components of the attention mechanism?

- ☒ a. Queries, Keys, and Values
- ☐ b. Keys, Weights, and Queries
- ☐ c. Inputs, Outputs, and Weights
- ☐ d. Scores, Values, and Gradients

**Question 5**

Complete

Mark 1.00 out of 1.00

Which of the following is a key step in the attention mechanism?

- ☐ a. Ignoring the context vector
- ☐ b. Selecting a random set of features
- ☐ c. Applying dropout to all layers
- ☒ d. Calculating the softmax values over all inputs

**Question 6**

Complete

Mark 1.00 out of 1.00

In the context of the attention mechanism, what does the "query" represent?

- ☐ a. The output of the decoder
- ☒ b. The representation of the task or input focus
- ☐ c. The weight matrix of the model
- ☐ d. The input sequence to the model

**Question 7**

Complete

Mark 1.00 out of 1.00

What is the primary purpose of the attention mechanism in neural networks?

- ☒ a. To focus on relevant parts of the input while processing data
- ☐ b. To improve activation functions like ReLU
- ☐ c. To increase the number of hidden layers
- ☐ d. To reduce the size of the dataset

**Question 8**

Complete

Mark 1.00 out of 1.00

In the tokenization process, what does `word_to_index[token] = len(word_to_index) + 1` accomplish?

- ☒ a. It maps unique tokens to their corresponding indices.
- ☐ b. It removes duplicates from the token list.
- ☐ c. It ensures tokens are indexed starting from 0.
- ☐ d. It initializes the word embedding layer.

**Question 9**

Complete

Mark 1.00 out of 1.00

Which of the following types of attention is commonly used in Transformer models?

- ☐ a. Local attention
- ☐ b. Temporal attention
- ☐ c. Global attention
- ☒ d. Multi-head self-attention

**Question 10**

Complete

Mark 1.00 out of 1.00

What does the `DataLoader` class achieve in the code?

- ☐ a. It computes the loss and gradients during backpropagation.
- ☒ b. It batches the data and allows shuffling for training.
- ☐ c. It preprocesses the text and sentiment labels.
- ☐ d. It splits the data into training and testing sets.